Policy gradient/ search is influential in NLP/ Proximal Policy Optimization (training GPT). The core idea:

1 Why Policy Gradient, not Value Based?

Intuition of Gradient Search

Approximate $V^\pi(s) \approx V_W(s)$ and $Q_w(s,a) \approx Q^\pi(s,a)$ by adjusting weight w.
Policy gradient: rather than generating policy from value (ϵ-greedy), directly parametrize policy with θ, i.e.
$$\pi_\theta(s,a) = \mathbb{P}[a|s;\theta]: \text{optimize } V(\theta) \text{ to find policy } \pi$$

The brief classification of policy gradient is as follows:

	Value-based	Policy-based	Actor-critic
Value function	learned	not present	learned
Policy	implicit (ϵ-greedy)	learned	learned

Instead of deterministic/ ϵ-greedy policies, need to focus heavily on **stochastic** for direct policy search!

- Repeated Trials, e.g. In rock paper scissors (of many rounds), deterministic policy is easily exploited by adversary.

- Boundary Condition, e.g. In gridworld, bound to only move one direction (else get stuck/ traverse for long time for slow convergence).

1.1 Gradient Free Policy Optimization

We begin with simple (but great) gradient-free baselines.

- Examples: Hill Climbing, Genetic Algo (evolution strategies, cross-entropy method, covariance matrix adaption)

- Known for decades but embarrassingly well: rivals standard RL techniques!

- Advantages: Flexible for any policy parameterization, easily to parallelize
  Disadvantage: Less sample efficient (ignores temporal structure)

2 Main Objective and Log-likelihood Trick for Policy Gradient

2.1 Policy Gradient

This section focuses on gradient descent; other popular algos include conjugate gradient and quasi-newton methods.
We assume **Episodic MDPs** for easy extension of objectives. We first outline the problem as follows:

Policy Objective Summary

- Goal: Given policy $\pi_\theta(s,a)$, find best parameter θ. Inherently, an optimization of $V(s_0, \theta)$ (i.e. the value fucntion depending on policy parameters).
- Purpose: Measure quality for policy π_θ with policy value at start state s_0.
- Works for: both episodic/ continuing and infinite horizons.

The method:

> **Vanilla Policy Gradient: Problem Formation**
>
> - Search the local maximum of policy value $V(s_0, \theta)$ with gradient increments:
>
> $$\Delta\theta = \alpha\nabla_\theta V(s_0, \theta) = \alpha \begin{pmatrix} \dfrac{\partial V(s_0, \theta)}{\partial\theta_1} \\ \vdots \\ \dfrac{\partial V(s_0, \theta)}{\partial\theta_n} \end{pmatrix}$$
>
> - Assumption: π_θ differentiable (and known gradient $\nabla_\theta\pi_\theta(s, a)$)
> - We can rewrite $V(s_0, \theta)$ in the following ways:
>   1. **Visited States and Actions**: $\mathbb{E}_{\pi_\theta}\left[\sum_{t=0}^{T} R(s_t, a_t); \pi_\theta, s_0\right]$
>   2. **Weighted Average of Q-values by Actions**: $\sum_a \pi_\theta(a|s_0)Q(s_0, a, \theta)$
>   3. **Trajectories Sampled using π_θ**: $\sum_\tau P(\tau|\theta)R(\tau)$

2.2 Log-Likelihood Trick and Score

In particular, it is of interest to consider writing $V(s_0, \theta)$ in trajectory form:

> To find the best policy parameter θ, we consider
>
> $$\arg\max_\theta V(\theta) = \arg\max_\theta \sum_\tau P(\tau; \theta)R(\tau)$$
>
> Taking gradient,
> $$\begin{aligned} \nabla_\theta V(\theta) =& \nabla_\theta \sum_\tau P(\tau; \theta)R(\tau) \\ & \sum_\tau \nabla_\theta P(\tau; \theta)R(\tau) (R \text{ being indep of } \theta) \\ & \sum_\tau \frac{P(\tau; \theta)}{P(\tau; \theta)}\nabla_\theta P(\tau; \theta)R(\tau) \\ & \sum_\tau R(\tau)P(\tau; \theta)\nabla_\theta \log P(\tau; \theta) \quad \text{(log-likelihood)} \end{aligned}$$
>
> **Approximate** in practice using m sample trajectories under π_θ:
>
> $$\nabla_\theta V(\theta) \approx \hat{g} = \frac{1}{m}\sum_{i=1}^{m} R(\tau^{(i)})\nabla_\theta \log P(\tau^{(i)}, \theta)$$

But trajectories can be decomposed into states and actions:

> $$\begin{aligned} \nabla_\theta \log P(\tau^{(i)}; \theta) =& \nabla_\theta \log\left[\mu(s_0)\prod_{t=0}^{T-1}\pi_\theta(a_t|s_t)P(s_{t+1}|a_{t+1}, s_{0:t}, a_{0:t})\right] \\ =& \sum_\tau \nabla_\theta \log \pi_\theta(a_t|s_t) \end{aligned}$$

Here

- We call $\sum_\tau \nabla_\theta \log \pi_\theta(a_t|s_t)$ the **score function**.

- the initial state $\mu(s_0)$ is constant; dynamics model $P(s_{t+1}|a_{t+1}, s_{0:t}, a_{0:t})$ is invariant to θ.

- In other words, no dynamics model is required to approximate the policy parameter θ.

Questions

1. Why trajectory form is practical ("better in training")?

   Two major reasons:
   (a) **Flexible**
       "Black box access" that only requires generated trajectory rollouts without differentiable envt. model.
       Allows turning $\nabla_\theta P(\tau)$ into $P(\tau)\nabla_\theta \log P(\tau)$, i.e. unbiased gradient estimate with simplicity.
   (b) **Easy Implementation**
       Transition probabilities $P(s_{t+1}|s_t, a_t)$ don't appear in gradient after log derivative.
       Only actions, episodes, states and rewards (not P) to compute gradient!

2. Why is log-likelihood ratio important here? What does it enable?

   Log trick enables PG without environment back-propagation.
   Without model **P** required, log likelihood enables additive (instead of multiplicative) decomposition.

2.3 Selecting a Right Policy

2.3.1 Softmax Policy

- In softmax, **exponentially weight** quantities of linear combination of features as probabilities (that add to 1):

$$\pi_\theta(s, a) = \frac{e^{\phi(s,a)^T \theta}}{\sum_a e^{\phi(s,a)^T \theta}}$$

- Then the score function can be written as $\nabla_\theta \log \pi_\theta(s, a) = \phi(s, a) - \mathbb{E}_{\pi_\theta[\phi(s,\cdot)]}$

2.3.2 Gaussian Policy

- A normal distribution is natural for continuous action spaces; at times used by deep NN.

- Action $a \sim N(\mu(s), \sigma^2)$. Mean $\mu(s) = \phi(s)^T \theta$ is a linear combination of state features.

- Then the score function can be written as $\nabla_\theta \log \pi_\theta(s, a) = \frac{(a - \mu(s))\phi(s)}{\sigma^2}$

What are the purposes of selecting Softmax and Gaussian policies?

Policy classes should allow (1) Easy action sampling and (2) Straightforward gradient $\nabla_\theta \log \pi_\theta(s, a)$ computation.
The two policies are both differentiable, allowing (2). For (1):
- **Softmax**: **Discrete** action probabilities that returns a nice (simple) gradient form.
- **Gaussian**: **Continuous** actions featuring straightforward gradient, common in robotics/ continuous control.

2.4 Summary: PG

A summary of policy gradient:

Intermediate Summary of PG

- Core idea: $\nabla_\theta V(\theta) = \mathbb{E}_{\pi_\theta}\left[\nabla_\theta \log \pi_\theta(s,a)Q^{\pi_\theta}(s,a)\right]$

Optimize $\arg\max_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta}\left[\sum_{t=0}^{\infty}\gamma^t r_t\right]$ with SGD on θ:

$$g = \nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta}\left[\sum_{t=0}^{\infty}\gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t)A^{\pi_\theta}(s_t,a_t)\right]$$

- State-action pairs with higher $\hat{Q}$ increases probabilities in average
- Direction of θ dependent on gradient of $\ln \pi(S_t, A_t, \theta)$ AND Q-values/ returns
- NOT guaranteed to converge to global optima (just local!)

When to use?
1. Differentiable reward functions; No dynamics required
2. Useful for both infinite horizon and episodic settings
3. Intuition: $R(\tau^{(i)})$ is replacable by other functions that *measures the wellness of sample x*
   Essentially, moving in the direction of $\hat{g}_i = f(x_i)\nabla_\theta \log p(x_i|\theta)$ pushes up the log probability proportionally.

The generalization of PG is as follows:

Policy Gradient Theorem

Assumption: Differentiable Policy $\pi_\theta(s,a)$, objective $J = \begin{cases} J_1 & \text{(Episodic)} \\ J_{avR} & \text{(Avg. Reward over time)} \\ \dfrac{1}{1-\gamma}J_{avV} & \text{(Avg. Value over time)} \end{cases}$ In any case, the policy

gradient is $\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}\left[\nabla_\theta \log \pi_\theta(s,a)Q^{\pi_\theta}(s,a)\right]$

Summary and Improvements of Policy-based RL

Criteria	Advantages	Disadvantages
Convergence	Better Properties	Typical Local Optimum
Flexibility	Effective in high-dim./ continuous action spaces	Inefficient, high-variance
	Can learn stochastic policies	policy evaluation

3 Variance Issues and Remedies

3.1 Problem: High Variance in PG

Currently, use

$$\nabla_\theta V(\theta) = \frac{1}{m}\sum_{i=1}^{m}R(\tau^{(i)})\sum_{t=0}^{T-1}\nabla_\theta \log \pi_\theta(a_t^{(i)}, s_t^{(i)})$$

to estimate

$$\nabla_\theta \mathbb{E}_\tau[R] = \mathbb{E}_\tau\left[\left(\sum_{t=0}^{T-1}r_t\right)\sum_{t=0}^{T-1}\nabla_\theta \log \pi_\theta(a_t|s_t)\right]$$

It is unbiased but noisy (high variance)!

On the high variance, several remedies serve as improvements:

3.2   Fix 1: Temporal Structure

- Focus on single reward item at once:

$$\nabla_\theta \mathbb{E}[r_{t'}] = \mathbb{E}[r_{t'} \sum_{t=0}^{t'} \nabla_\theta \log \pi_\theta(a_t|s_t)]$$

$$V(\theta) = \nabla_\theta \mathbb{E}[R] = \mathbb{E}[\sum_{t'=0}^{T-1} r_{t'} \sum_{t=0}^{t'} \nabla_\theta \log \pi_\theta(a_t|s_t)]$$

- Sum up over t to obtain:

$$= \mathbb{E}[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t, s_t) \sum_{t'=t}^{T-1} r_{t'}] \quad \text{(because later decisions don't influence past rewards}$$

- This can be further simplified: trajectory $\tau^{(i)}$ has return $G_t^{(i)} = \sum_{t'=t}^{T-1} r_{t'}^{(i)}$. Hence

$$\nabla_\theta \mathbb{E}[R] \approx \frac{1}{m} \sum_{i=1}^{m} \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t, s_t) G_t^{(i)}$$

Monte-Carlo Policy Gradient

Making use of likelihood ratio / score function and temporal structure, update param θ with

$$\Delta \theta_t = \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) G_t$$

after initializing θ arbitrarily.

Q: How does Temporal structure reduce variance?

- Each $r_{t'}$ is paired with actions ONLY from time 0 to t'.
- Actions after t' does NOT affect $r_{t'}$:
  Assigning rewards to ONLY actions that influence it achieves noise reduction.

3.3   Fix 2: Baseline Function

As iteration costs time and computational resources, we desire quick convergence to local optima.

Baselines: Unbiasedness and Other Considerations

$$\nabla_\theta \mathbb{E}_\tau[R] = \mathbb{E}_\tau \left[\left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t) \right]$$

Why it works?
- Unbiased for any b if b is a function of s but not θ, because
  $\mathbb{E}_\tau [\nabla_\theta \log \pi(a_t|s_t; \theta) b(s_t)] = 0$

- A near-optimal baseline choice is the expected return $b(s_t) \approx \mathbb{E} \left[\sum_{t'=t}^{T-1} r_{t'} \right]$

- Other choices: State-value function $V^\pi(s) = \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)]$

Interpretation: Increase logprob of action at proportionally to how much returns $\sum_{t'=t}^{T-1} r_{t'}$ are better than expected.

Mathematically, baseline functions achieve variance reduction as follows:

Core idea: Break down t in the summation:

$$\text{Var}[\nabla_\theta[R]] = \text{Var}[\sum_{t=0}^{T-1} \nabla_\theta \log \pi(a_t|s_t;\theta)(R_t(s_t) - b(s_t))]$$

- As we sample from trajectories τ (MC),

$$\approx \sum_{t=0}^{T-1} \mathbb{E}_\tau \left[\text{Var}[\nabla_\theta \log \pi(a_t|s_t;\theta)(R_t(s_t) - b(s_t))]\right]$$

- For each t, write variance $\text{Var}[X]$ as $\mathbb{E}[X^2] - (\mathbb{E}[X])^2$:

$$\mathbb{E}\left[((\nabla_\theta \log \pi(a_t|s_t;\theta))(R_t(s_t) - b(s_t)))^2 \right] - \mathbb{E}\left[((\nabla_\theta \log \pi(a_t|s_t;\theta))(R_t(s_t) - b(s_t))) \right]^2$$

- Second term is not affected by choice of $b(s)$ (unbiased = same expectation).
  The variance equals $\underset{b}{\arg\min} \quad \mathbb{E}\left[((\nabla_\theta \log \pi(a_t|s_t;\theta)))^2 ((G_t(s_t) - b(s_t)))^2 \right]$

  $$= \underset{b}{\arg\min} \quad \mathbb{E}_{s \sim d^\pi} \left[\mathbb{E}_{a \sim \pi(\cdot|s), G|s,a} \left[((\nabla_\theta \log \pi(a_t|s;\theta)))^2 ((G_t(s_t) - b(s_t)))^2 \right] \right]$$

- A **weighted least-squares** problem that minimizes

$$\sum_i \sum_t |b(s_t^i) - G_t^i|^2 \qquad (G \text{ (or } A = G - b) \text{ being target)}$$

  with solution (after taking zero gradient)

$$b(s) \approx \mathbb{E}_{a \sim \pi(\cdot|s), G|s,a}[G_t(s)]$$

Q: Intuitively, what are the uses of baseline functions? How does variance get reduced?

- Second term: If $b(s)$ is "close" to true $V^\pi(s)$ (high correlation between the two), $G_t(s_t) - b(s_t)$ achieves lower variance.

- First term: Log-probability only updates strongly (decisively, to increase variance) if significant return differences to baseline is observed.

Q: What does it mean by "Near optimal"? When is it not?

- According to above, when variance reduction of the second term > variance increase of first term.
  Usually the case as $b(s) \approx V^\pi(s)$ typically, achieving maximal variance reduction in theory.

- If learned baseline is inaccurate (i.e. poor function approximator/ predictor), less reduction.
  Extreme cases: Taking $b(s) = $ const or 0 could be better.

3.4 Fix 3: Actor-Critic

The original G_t^i estimates expected discounted sum of returns (from single roll): unbiased but high variance. To solve this:

- Leverage bootstrapping and approximation (similar to TD vs MC and VFA) to introduce bias.

- Use "Critic" to estimate the ratio $\dfrac{V}{Q}$. The popular class of "Actor-critic" methods explicitly represents (and updates) policy and values.

- Essentially, replace $\sum_{t'=t}^{T-1} r_{t'} - b(s_t)$ [Vanilla MC] with Q-values $(Q(s_t, a_t; \mathbf{w}) - b(s_t))$ [This is essentially TD]
  or advantage function $\hat{A}^\pi(s_t, a_t)$ where $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$.

Alternative Targets to MC Estimators

With vanilla MC, the gradient is estimated by

$$\nabla_\theta V(\theta) \approx \frac{1}{m} \sum_{i=1}^{m} \sum_{t=0}^{T-1} R_t^i \nabla_\theta \log \pi_\theta(a_t^{(i)}|s_t^{(i)})$$

Proposed replacements:
1. N-step estimators:
   - $\hat{R}_t^{(1)} = r_t + \gamma V(s_{t+1})$, $\hat{R}_t^{(2)} = r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2})$
   - $\hat{R}_t^{(\infty)} = r_t + \gamma r_{t+1} + \gamma r_{t+2} + \dots$
2. Advantage estimators, by *subtracting baselines of $V(s_t)$ from above*
   - $\hat{A}_t^{(1)} = \hat{R}_t^{(1)} - V(s_t) = r_t + \gamma V(s_{t+1}) - V(s_t)$ (*low variance, high bias*)
   - $\hat{R}_t^{(\infty)} - V(s_t) = r_t + \gamma r_{t+1} + \gamma r_{t+2} + \dots - V(s_t)$ (*high variance, low bias*)

What do the words "actor" and "critic" mean?

- "actor": Policy π_θ selecting actions
- "critic": VFA (say, trained via temporal-difference) to "criticize" actions chosen by actor
- Proposed $Q^\pi(s,a)$ or $A^\pi(s,a)$ serve as baselines/ targets for actor updates to reduce variance.

4 PPO and Its Two Variants

4.1 Problem: Poor Sample Efficiency of PG

A PG algo should minimize # iterations to reach a good (probably suboptimal) policy within time. The limitations of vanilla PG:

4.1.1 Poor sample efficiency

Variance reduces slowly (even after improvements), because PG is an **on-policy** expectation: Data immediately discarded after just one gradient step

- Collect sample estimates from trajectories of same policy (more stable), or other policies (off-policy, less stable).

- Opportunity: Can we take multiple gradient steps from old data before new policy?

Problems of Determining Gradient Step

Problem: Difficult to handle step size (dist. in parameter space $\neq$ dist. in policy space)
- e.g. Matrices in tabular case $\Pi = \{\pi : \pi \in \mathbb{R}^{|S| \times |A|}, \sum_a \pi_{s_a} = 1, \pi_{s_a} \geq 0\}$

  VS steps of policy gradient in parameter space $\implies$ unable to map/ gauge size!
- SGD of $\theta_{k+1} = \theta_k + \alpha_k \hat{g}_k$ is subject to performance collapse with large steps!
  e.g. logistic function: small $\Delta\theta$ leads to big policy changes

4.2 Policy Performance Bounds

The solution to restrict policy change from more than intended is through policy performance bounds:

Distance in Value to Policy

To respect distance mapping in policy space, exploit relationships between policy performance:

$$J(\pi') - J(\pi) = \mathbb{E}_{\tau \sim \pi'}\left[\sum_{t=0}^{\infty} \gamma^t A^\pi(s_t, a_t)\right] = \frac{1}{1-\gamma}\mathbb{E}_{s \sim d^{\pi'}, a \sim \pi'}[A^\pi(s,a)]$$

Here $d^\pi(s) = (1-\gamma)\sum_{t=0}^{\infty} \gamma^t P(s_t = s | \pi)$ is the **weighted** distribution of states. Making use,

$$\max_{\pi'} J(\pi') = \max_{\pi'} J(\pi') - J(\pi) \qquad\qquad = \max_{\pi'} \mathbb{E}_{\tau \sim \pi'}\left[\sum_{t=0}^{\infty} \gamma^t A^\pi(s_t, a_t)\right]$$

Now, rewrite the objective:

$$= \max_{\pi'} \frac{1}{1-\gamma}\mathbb{E}_{s \sim d^{\pi'}, a \sim \pi'}[A^\pi(s,a)]$$

$$= \max_{\pi'} \frac{1}{1-\gamma}\mathbb{E}_{s \sim d^{\pi'}, a \sim \pi}\left[\frac{\pi'(a|s)}{\pi(a|s)} A^\pi(s,a)\right]$$

Why rewrite?

1. Now, performance of π' is defined in advantages from π.

2. Requires trajectories sampled from π' (desired: from π because our tweak features $s \sim d^\pi$).

We have a useful approximation:

Relative Policy Performance Bounds

$$J(\pi') - J(\pi) \approx \mathbb{L}_\pi(\pi') \qquad \text{for close } \pi' \text{ and } \pi \ (d^{\pi'} = d^\pi)$$

Approximation quality is ensured by relative policy performance bounds:

$$|J(\pi') - (J(\pi) + \mathbb{L}_\pi(\pi'))| \leq C\sqrt{\mathbb{E}_{s \sim d^\pi}[D_{KL}(\pi'||\pi)[s]]}$$

But what is $\mathbb{E}_{s \sim d^\pi}[D_{KL}(\pi'||\pi)[s]]$?

KL-Divergence

Such divergence measures distance between probability distributions:

$$D_{KL}(P||Q) = \sum_x P(x) \log \frac{P(x)}{Q(x)}$$

KL satisfies $D_{KL}(P||P) = 0, D_{KL}(P||Q) \geq 0, D_{KL}(P||Q) \neq D_{KL}(Q||P)$.
Between policies, $D_{KL}(\pi'||\pi)[s] = \sum_{a \in A} \pi'(a|s) \log \frac{\pi'(a|s)}{\pi(a|s)}$

After the approximation, we can optimize using trajectories sampled from the old policy π!

Policy Optimization under Bounded KL Approximation

Policy improvement can be estimated by sampling from old policy π!

$$J(\pi') - J(\pi) \approx L_\pi(\pi') = \mathbb{E}_{\tau \sim \pi}\left[\sum_{t=0}^{\infty} \gamma^t \frac{\pi'(a_t, s_t)}{\pi(a_t, s_t)} A^\pi(s_t, a_t)\right]$$

4.3   Variants of PPO

Taking advantage of policy performance bounds, PPO penalizes large policy change iterations. There are two methods:
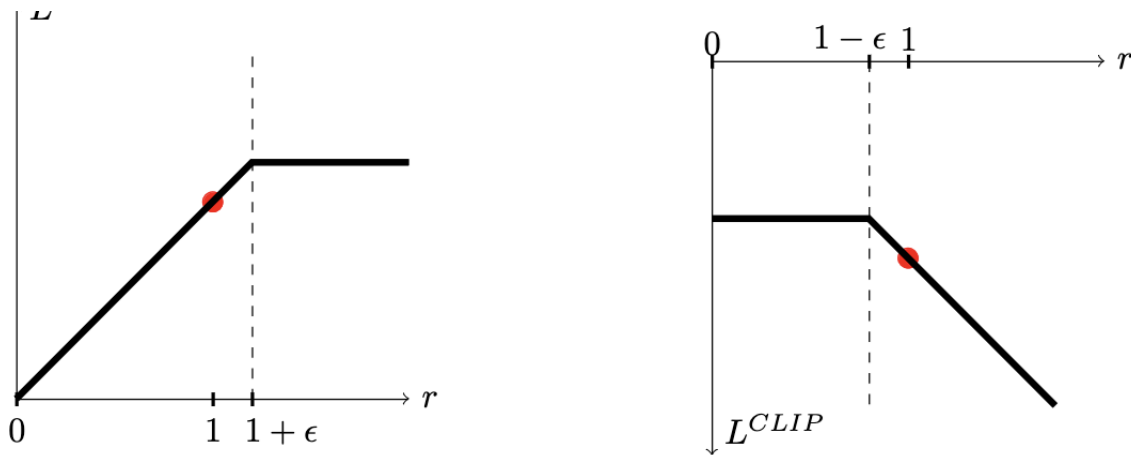
1. regularization term on KL-divergence (Adaptive Penalty)

2. pessimistic objective on far-away policies (Clipped Objective)

Adaptive Penalty

$$\theta_{k+1} = \arg\max_{\theta} L_{\theta_k}(\theta) - \beta \bar{D}_{KL}(\theta||\theta_k)$$

Here, KL-divergence is an expectation:

$$\bar{D}_{KL}(\theta||\theta_k) = \mathbb{E}_{s \ sim d^{\pi_k}} D_{KL}(\theta_k(\cdot|s), \pi_\theta(\cdot|s))$$

Penalty coefficient β_k changes between iterations:
- Initiate policy param θ_0, initial KL penalty β_0, target KL-divergence δ
- Compute policy update (iterate θ) by K steps of minibatch SGD (via Adam)
- Control KL-divergence to be around δ by adjusting penalty:
  If $\bar{D}_{KL}(\theta||\theta_k) \geq 1.5\delta$ then $\beta_{k+1} = 2\beta_k$;
  elif $\bar{D}_{KL}(\theta||\theta_k) \leq \dfrac{\delta}{1.5}$ then $\beta_{k+1} = \dfrac{1}{2}\beta_k$

This KL penalty is called "adaptive" because of how β_k changes (adapts quickly) according to KL-divergence.

An alternative approach is clipping: restrict via pessimistically treating objective value far away from θ_k.

Clipped Objective on Policy Changes

- Define relative probability change: $r_t(\theta) = \dfrac{\pi_\theta(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)}$
- The new objective is

$$L_{\theta_k}^{\text{Clip}}(\theta) = \mathbb{E}_{\tau \sim \pi_k}\left[\sum_{t=0}^{T}\left[\min\left(r_t(\theta)\hat{A}_t^{\pi_k}, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon)\hat{A}_t^{\pi_k}\right)\right]\right]$$

  In other words, $r_t(\theta)$ is clipped between $(1-\epsilon, 1+\epsilon)$; hyperparameter ϵ usually set at 0.2.
- Policy update: $\theta_{k+1} = \arg\max\quad_\theta L_{\theta_k}^{\text{Clip}}(\theta)$

Here, Clipping discentivizes going far from θ_{k+1}:

- Left graph below: when the advantage function $A > 0$; right graph when $A < 0$.

- $L_{\theta_k}^{\text{Clip}}(\theta)$'s increase is suppressed at extreme values towards the sign of A.

- Clipping is simple to implement but works well compared to KL penalty.

PPO's performance consistently tops other algos, hence wildly popular.

- Today, it is a key component of ChatGPT (readings: OpenAI blog (2017), Publication by Schulman et. al. (2017))

- Different outcomes with reward scaling/ learning rate annealing.

Why KL-divergence applied is an expectation?

- Recall that the two distributions (of the same variable) occurs in probabilities.

- To compare, take average (expectation) over states (discounted distribution).

How does the two methods compare?

KL Penalty
- Idea: Tune penalty β adaptively to control KL around the target threshold

- Pros: Conceptually direct (literally penalize KL)

- Cons: Tricky to schedule/ tune/ stabilize β

Clipped Objective
- Idea: Force objective $r_t\theta$ to plateau when significantly deviate from 1.

- Pros: Simple to implement and good practical performance

- Cons: Less direct than penalty/ constraint

In practice, clipping is much more common/ popular given implementation and performance.