

1 Introduction

1.1 Overview

1.1.1 On General Machine Learning

ML has seen a huge (exponential) increase in academic interest since 2000. It is categorized as follows:

1. Supervised and unsupervised Learning (SL/ UL): Only involves generalization \implies Differs in whether there is labelling
2. Imitation Learning (IL): behaviour cloning, assuming input demonstrations of GOOD policies
3. **Reinforcement Learning (RL): model-based learning (i.e. given reward info, states reached and actions taken)**

1.1.2 On Reinforcement Learning

RL refers to **Sequential decision making** to make good decisions under uncertainty.

Steps of RL

RL involves the following four steps:

1. **Optimization:** Find optimal decisions yielding best/ very good outcomes (e.g. minimum distance inter-city route).
2. **Delayed consequences:** Plan and reason long term ramifications (e.g. saving for retirement), through temporal credit assignment (i.e. identifying which past actions led to current rewards?).
3. **Exploration:** Learn/ explore by making decisions (that impacts/ reinforces what we learn about from rewards, e.g. riding and falling from a bike).
4. **Generalization:** Policy maps (= compresses) past experience to action.

Why not always choose the best-known action in decision making?

New actions might lead to even better rewards (exploration-exploitation trade-off).

Example: Trying a new restaurant vs. Going to favorite one?

1.1.3 Course Flow

- High-level learning goals: Understand theoretical and empirical approaches for evaluating reinforcement learning algorithm quality
- Flow: First explore MDP \implies model-free (policy evaluation and control) \implies function approximation \implies policy search + exploration

1.2 Basic Layout of RL

RL algorithms involve **State, actions, reward model and dynamics model**:

- RL aims to select sequence of actions to maximize $E[\text{future reward}]$,
- hence involves balancing immediate and long term rewards.

Example: Choose sequence of web advertisement to maximize view time

Mathematical Formulation of RL algorithms

For each time t , a reinforcement learning agent (A) that interacts with the world (W) would...

1. A takes **action** a_t ,
2. W updates a_t , emits **observation** o_t and **reward** r_t
3. A receives **history** h_t containing a_i, o_i, r_i for all time i , cumulative history is called state s_t .

In such sequential dynamic programming problems of RL, we have the following considerations:

1. **States**: Is the state Markov? Is the world partially observable?
2. **Dynamics**: Are dynamics deterministic or stochastic?
3. **Horizon of effect**: is future states taken into account?

1.2.1 States: Markov Assumption

In many cases, we have the Markov assumption:

Markov Assumption: Definition

The future is independent of past given present (i.e. **state** is a sufficient statistic of **history**):

$$P(s_{t+1}|s_t, a_t) = P(s_{t+1}|h_t, a_t)$$

1.2.2 Dynamics: Model, Policy, Value Function

RL Algos must contain ≥ 1 of Model, Policy, Value Function:

Model

A MDP (Markov Decision Process) contains

1. **Transitional dynamics model**, predicting next agent state $P(s_{t+1} = s' | s_t = s, a_t = a)$
2. **Reward model**, predicting immediate reward $r(s_t = s | a_t = a) = E[r_t | s_t = s, a_t = a]$

Why Markov?

- Simple, as long as history is in part of state (in practice, assume $s_t = o_t$)
- s_t affects computational complexity, data required and result performance.

Why do we start with MDP?

Structured framework to model decision-making in uncertain environments.

\implies Easier to apply RL algorithms, defining **states, actions, rewards, and transitions**.

Policy

Policy (denoted by π) determines how actions are chosen (i.e. $\pi : S \rightarrow A$). It can be

1. **Deterministic**: $\pi(s) = a$, or
2. **Stochastic**: $\pi(a|s) = P(a_t = a | s_t = s)$

Value Function

Value function (denoted by V^π) discounts sum of future rewards under policy π , to quantify goodness of states/ actions:

$$V^\pi(s_t = s) = \mathbb{E}_\pi [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s]$$

Discounting Factor Considerations

Discount factor γ weighs immediate vs future rewards:

- $\gamma = 0$: only cares about immediate;
- $\gamma = 1$: possible for finite episode length.
- To prevent infinite returns, conveniently set $\gamma \leq 1$.

1.3 Agent and Feedback: MRP

RL agents can be either

1. **Model-based**: explicit model, may (not) have policy and value functions;
2. **Model-free**: no model, but value and/or policy functions must be explicit.

Learning Mechanisms

Given a model of world (dynamics and reward), of finite set of states and actions.

Two actions for the RL agent include:

1. **Evaluation**: Estimate/ predict expected rewards given policy.
2. **Control**: find the best policy via optimization.

Now consider the learning algorithm under Markov Chain (Markov Process):

- Markov processes involve sequences of random states with Markov property
They are defined by (finite) set of states S and transition model P , specifying $P(s_{t+1} = s' | s_t = s)$
- In a transition model, there is **no reward**, hence no actions are defined.
For finite (N) states, express P as $N \times N$ matrix.
- To bridge the stochasticity in quantifying rewards/ value, Markov Reward Processes (MRP) and Markov Decision Processes (MDP) are used.

1.3.1 Markov Reward Processes (MRP)

MRP **includes the reward** into consideration.

- It requires states S , transition model $P_{n \times n}$, reward $R(s_t = s)$, and discount factor $\gamma \in [0, 1]$.
- With H steps in each episode (can be infinite), the return

$$G_t = r_t + \gamma r_{t+1} + \dots + \gamma^{H-1} r_{t+H-1}$$

- Define the state value $V(s) = E[G_t | S_t = s] = R(s) + \gamma \sum_{s' \in S} P(s' | s) V(s')$.

The first term is the immediate reward; the second term is the discounted sum of future rewards.

Obviously we are interested in computing the value of a MRP. There are two methods:

Analytical and Iterative Solutions of MRP

Method 1 (Matrix inverse) involves writing in matrix-vector forms:

- If the number of states n is finite, rewards \mathbf{R} is a vector.
- Express $V(s)$ in the MRP using a matrix equation

$$\mathbf{V}_{n \times 1} = \mathbf{R}_{n \times 1} + \gamma \mathbf{P}_{n \times n} \mathbf{V}_{n \times 1}$$

- Rearranging, the solution $\mathbf{V} = (\mathbf{I} - \gamma \mathbf{P})^{-1} \mathbf{R}$.
Direct solve takes around $O(n^3)$ time.

Method 2 (Iterative Solution) involves looping until convergence:

- Dynamic programming by initializing $V_0(s) = 0$ for all s
- For each iteration k , and all states $s \in S$:
$$V_k(s) = R(s) + \gamma \sum_{s' \in S} \mathbb{P}(s'|s) V_{k-1}(s')$$
- Each iteration takes $O(n^2)$ time (because $n = |S|$).

Why $\mathbf{I} - \gamma \mathbf{P}$ is Invertible?

With $\gamma < 1$ (ensuring convergence), and the fact that \mathbf{P} is stochastic (rows sum to 1):

- The eigenvalues of P are at most 1.
- When scaled by $\gamma < 1$, largest eigenvalue $\lambda_1 < \gamma = 1$.
This ensures $\mathbf{I} - \gamma \mathbf{P}$ is full rank (diagonally dominant).

2 How to make good decisions given MDP?

2.1 Solving MDP by reducing to MRP

MDP as MRP Expansion

- $\text{MRP}(S, P, R, \gamma) + \text{Actions } A = \text{MDP}(S, A, P, R, \gamma)$

Example: Consider a robot vacuum moving around a room:

- Only models probabilities of movement without considering actions: MRP
- Includes actions like 'turn left' or 'move forward': MDP

- $\text{MDP}(S, A, P, R, \gamma) + \text{Policy } \pi(a|s) = \text{MRP}(S, P^\pi, R^\pi, \gamma)$
Use same technique of computing MRP to evaluate MDP policy!
- Various settings of MDPs:
 1. **Single-state:** Bandits
 2. **Continuous states:** Optimal control
 3. **If state is history:** Partially observable MDPs (POMDPs)

We can solve MDP as an MRP as follows:

Solving MDP Iteratively

For $k = 1$ till convergence, for any $s \in S$, the state value function V is iterated with:

$$V_k^\pi(s) = \sum_a \pi(a|s)[r(s|a) + \gamma \sum_{s' \in S} P(s'|s, a)V_{k-1}^\pi(s')] \quad \text{where } s' \in S$$

This is called a Bellman backup equation.

- Current value at k th iteration is obtained by **summing the policy probabilities** varying state s , **weighted by the value** (current reward + future rewards discounted by one step).
- For deterministic $\pi(s)$, i.e. $P(\pi(s)|s) = 1$, the formula reduces to

$$V_k^\pi(s) = r(s|\pi(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi(s))V_{k-1}^\pi(s')$$

Where does Bellman Equation arise?

It arises naturally when breaking down long-term decision problems (by recursion) into smaller subproblems, as in policy evaluation.

2.2 How to find the optimal policy under MDP?

- Given $|A|$ distinct options for each of the $|S|$ states, there are a total of $|A|^{|S|}$ deterministic policies possible.
- Optimal policy may NOT be unique (but is deterministic and stationary (does not depend on time step) for MDP in infinite horizon!), but there exists a **unique optimal value function**.
- Such optimal policy can be found by "controlling", i.e. optimizing $\pi(s) = \arg \min_{\pi} V_{\pi}(s)$

Why is such optimal policy deterministic?

- Randomness in an MDP **comes from state transitions**, not from the policy itself.
- Given enough iterations, an optimal policy learns **best action in every state** *implies* randomness is unnecessary.

There are primarily **THREE methods** in finding the optimal policy under MDP:

2.2.1 Policy iteration (PI)

PI is more efficient than enumeration. Its process is as follows:

PI Process

- (1) **Initialization**: Initialize $\pi_0(s)$ randomly for all states s .
- (2) **Check L1 Norm in Loop**: Iterate when **i == 0** OR $\|\pi_i - \pi_{i-1}\|_1 > 0$,
i.e. *if policy changes for any state*.
Each time, compute $V_{\pi}(i)$ (policy evaluation), and π_{i+1} (policy improvement) to iterate new **i**

But why does PI guarantee generating better policies?

- Define the **state action value** (Q-function):

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a)V_{\pi}(s')$$

It differs from V_{π} that **action a is taken in addition** to following the policy.

- Then for new policy π_{i+1} ,

$$\pi_{i+1}(s) = \arg \min_a Q^{\pi_i} \pi(s, a) \forall s \in S$$

is a deterministic quantity.

- To improve the state-action value by varying a ,

$$\max_a Q^{\pi_i}(s, a) \geq Q^{\pi_i}(s, \pi_i(s)) = R(s, \pi(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi_i(s)) V^{\pi_i}(s') = V^{\pi_i}(s)$$

This is more rigorously defined as follows:

Monotonic Improvement in Policy

Definition: $V^{\pi_1} \geq V^{\pi_2}$ if $V^{\pi_1}(s) \geq V^{\pi_2}(s)$ for all states s . (Strict inequality $V^{\pi_{i+1}} > V^{\pi_i}$ if π_i is suboptimal.)

Implication: By taking $\pi_{i+1}(s)$ for one action and following π_{i+1} forever, expected sum of rewards is at least as good as always following π_i . Since we always pick a better action, policies never get worse!

Regarding change (and iteration) of policies in PI

- If policy doesn't change, it cannot be changed anymore.
- There is a maximum number of policy iterations (since a finite number of options for each finite state).

Why state action value is defined? Why store Q instead of V?

- It evaluates the expected return taking action a first, then act optimally afterward.
- Instead of storing $V(s)$, store $Q(s, a)$ to directly optimize actions.

2.2.2 Value iteration (VI)

The idea is to maintain **optimal value** of starting in state s , if finite number of steps k left in episode.

Then, iterate to consider longer and longer episodes:

Bellman Operator

For a value function, the Bellman Operator B returns a new value that improves if possible:

$$BV(s) = \max_a \left[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V(s') \right]$$

For a policy,

$$B^\pi V(s) = R^\pi(s) + \gamma \sum_{s' \in S} P^\pi(s'|s) V(s')$$

Example: Updating est. travel time in Google Maps as new traffic data arrives (i.e. iteratively refines value estimates).

How to "maintain optimal value"? Why does a disappear?

- Maintaining optimal value means ensuring that for every state, compute best expected return by considering all possible actions.
- As we only keep the highest value across actions, we remove a .
- We only care about values during iteration: Extract optimal policy by checking which action led to that value

VI Process

- (1) **Initialization:** $k = 1$, $V_0(s) = 0$ for all states s .
- (2) **Check Norm in Loop:** Loop until convergence (i.e. checking $L_\infty = \|V_{k+1} - V_k\|_\infty \leq \epsilon$)

$$\text{For each state } s \text{ in } S, \text{ iterate } V_{k+1}(s) = \max_a \left[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') \right]$$

This iteration step is equivalently $V_{k+1} = BV_k$

For the optimal policy for **finite horizon** H , iterate (loop) from $i = 0 : H$ (exclusive of H):

- $V_{k+1}(s) = \max_a \left[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') \right]$
- $\pi_{k+1}(s) = \arg \max_a \left[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') \right]$

But why does VI guarantee convergence?

- Quantify convergence with **contraction operator** O , satisfying

$$\|OV - OV'\|_n \leq \|V - V'\|_n \text{ for any } L_n \text{ norm.}$$

- Converges when $\gamma < 1$ OR ending up in terminal state with probability 1!
- VI converges to **unique solution** for discrete state and action spaces
- Policy evaluation = compute **fixed point** of B_π , by repeatedly applying operator until V stops changing, i.e.

$$V^\pi = B^\pi B^\pi \dots B^\pi V$$

What does “fixed point” here mean?

- Applying the operator doesn't change the result anymore.
- Signifies convergence (e.g. in PI, determine whether the policy should be updated).

Does the initialization of values in value iteration impact anything?

- It only affects convergence speed, but same final result.
- Because Bellman backup is a contraction mapping.

A brief summary of PI vs VI

Type	PI	VI
To compute...	Infinite horizon value	Optimal finite k - horizon value
Use	Select another (better) policy	Find optimal policy, by incrementing time horizon
Convergence	#Iterations \leq Enumeration of all deterministic policies, i.e. of size $ A ^{ S }$	Depends on discounting factor γ , possible to exceed $ A ^{ S }$

2.3 Monte Carlo (MC) Methods for MDP

Given dynamics and reward models, can do policy evaluation through DP.

$$V_k^\pi(s) = r(s|\pi(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi(s)) V_{k-1}^\pi(s')$$

Here, V_k^π is an estimate of V^π , and **the second part** can be replaced by **bootstrapping** if dynamics/ reward models are unknown but deterministic : when there is **no access to true MDP**.

- MC follows policy π to **generate sample trajectories** and take expectation on sample return to approximate (estimate) $V^\pi(s)$.
- Rationale: Mean return converges to value (by LLN)
Note: All trajectories may not be of the same length because of MDP.
- Requires **episodic settings (terminal state/ finite horizon)** to end episode
 \Rightarrow Terminate episode before averaging.

Advantages of MC

- Weak assumption (No need Markov assumption/ knowing MDP dynamics and rewards)
- Sometimes preferred over DP for policy evaluation even with known dynamics model and reward. (Explain why?)

The process and variations of MC is as follows:

MC Process

1. **Initialize:** Number of samples $N(s) = 0$, grand total $G(s) = 0$ for each state $s \in S$.
2. **Loop:**
 - For each episode i , sample trajectories $(s_{i,1}, a_{i,1}, r_{i,1}), (s_{i,2}, a_{i,2}, r_{i,2}), \dots, s_{i,t_i}$
 - Compute return value $G_{i,t}$ as return from time step t :
$$G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \dots + \gamma^{T_i-t} r_{i,T_i}$$
 - Update state values for each time step t until end of episode i :
Increment # visits: $N(s) = N(s) + 1$
Increment total return: $G(s) = G(s) + G_{i,t}$
Update estimate $V^\pi(s) = G(s)/N(s)$

There are **Three variations** for MC:

Variations of MC

- (i) **First-time MC:** Increment $N(s)$ for the **first time visiting state s** in episode i , i.e. $N(s) = 1$
- (ii) **Every-time MC:** Increment $N(s)$ **each time visiting state s** in episode i
- (iii) **Incremental MC:** After obtaining the new return value $G_{i,t}$, rewrite the update from $V^\pi(s)$ to $V^\pi(s)'$:

$$\begin{aligned}
 V^\pi(s)' &= \frac{G(s)'}{N(s)'} = \frac{G(s) + G_{i,t}}{N(s) + 1} = \frac{G(s)}{N(s) + 1} + \frac{G_{i,t}}{N(s) + 1} \\
 &= V^\pi(s) \times \frac{N(s)}{N(s) + 1} + \frac{G_{i,t}}{N(s) + 1} \\
 &= V^\pi(s) + \frac{1}{N(s) + 1} (G_{i,t} - V^\pi(s)) \\
 V^\pi(s)' &= V^\pi(s) + \alpha (G_{i,t} - V^\pi(s))
 \end{aligned}$$

The general form replaces the learning rate $\alpha = 1/[N(s) + 1]$ to other rates. Viewed from this perspective:

MC Type	Every Visit	First Visit	General
α	$\frac{1}{N(s_{i,t})}$	1 if $N(s) = 0$, 0 otherwise	$V^\pi(s)$
Properties	Biased, but consistent and offers better MSE	Unbiased estimator of $V^\pi(s)$ by LLN	Rate larger than $1/N(s_{i,t})$ is

Q: What is an intuitive explanation to such every-visit "bias"? What does it mean by helpful in non-stationary domains?

2.4 Evaluating Policy Estimation Quality

We take theoretical and computational aspects for consideration:

- Theoretical: Statistical efficiency, consistency and empirical accuracy (MSE)
- Computation: Runtime and Memory Complexity

2.4.1 Bias and Variance

The empirical accuracy of learning models is quantified by Mean Squared Error (MSE):

$$\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta}) + \text{Bias}(\hat{\theta})^2$$

- $\text{Var}(\hat{\theta}) = \mathbb{E}_{x|\theta}[(\hat{\theta} - \mathbb{E}[\hat{\theta}])^2]$
- $\text{Bias}(\hat{\theta})^2 = (\mathbb{E}_{x|\theta}[\hat{\theta}] - \theta)^2$

Consistency quantifies how reliable an estimator is:

- If there are n data points of x , $\lim_{n \rightarrow \infty} \mathbb{P}(|\hat{\theta}_n - \theta| > \epsilon) = 0$ for arbitrary ϵ .
- In other words, estimates (produced by the estimator) "converge" to the true parameter value.

An unbiased estimator may not be consistent. Two simple examples:

- The variance estimator $\frac{1}{n} \sum (y_i - y_n)^2$ is biased but consistent.
- The data point X_1 as an estimator of μ for $N(\mu, \sigma^2)$ is unbiased but NOT consistent.

2.4.2 How does MC fare under those criteria?

(1) On Convergence: For MC, under varying learning rate α , incremental MC will converge to true value of policy, $V^\pi(s_j)$ if

1. $\sum_{n=1}^{\infty} \alpha_n(s_j) = \infty$, but
2. $\sum_{n=1}^{\infty} \alpha_n^2(s_j) < \infty$.

(2) On MSE: MC yields a **high-variance estimator** that takes a lot of data to reduce.

3 Value Function Approximation

Recall in optimizing policy π , we can either **evaluate** (compute Q^π) or **improve** (update π given Q^π).

Problem: Deterministic policies in approximation yield **no exploration!**

- If π deterministic, $Q(s, a)$ is undefined for $a \neq \pi(s)$. Can't learn about actions without trying them!
- Exploitation of past experience that proved high reward.

Solution: Assign probability on unvisited pathways!

- ϵ -greedy pathways selects argmax action with **$1-\epsilon$ probability**, else random. i.e. $\pi(a|s) = \begin{cases} \arg \max_a Q(s, a), & \text{probability } 1-\epsilon \\ a' \neq \arg \max_a Q(s, a), & \text{probability } \epsilon \end{cases}$
- ϵ -greedy policy guarantees monotonic improvement, and can estimate $Q(s, a) \forall s, a$.

3.0.1 Greedy in Limit of Infinite Exploration (GLIE)

GLIE: If $\lim_{i \rightarrow \infty} N_i(s, a) \rightarrow \infty$ (i.e. all state-action pairs visited infinite times)

- Behavior policy (policy used to act in the world) converges to greedy policy
- GLIE is ϵ -greedy if $\epsilon = \frac{1}{i} \rightarrow 0$.
- MC under GLIE converges to optimal value, i.e. $Q(s, a) \rightarrow Q^*(s, a)$

Q: What is the convergence to optimal Q^* function and the expected performance of MC control with ϵ -greedy policy? When may MC online control fail to compute the optimal Q^* ?

Other than on-policy learning (through direct experience following the policy), **off-policy** learning exists (experience gathered from [following a different policy](#)).

In particular, Q-learning is a popular area of off-policy learning:

Q-learning Definition and Properties

- Q-learning estimates Q^{π^*} by acting with another π_b by maintaining Q estimates and bootstrapping.
- Update step:

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha \left([r_t + \gamma \arg \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t)] \right)$$

where $\arg \max_{a'} Q(s_{t+1}, a')$ is by following another behaviour policy.

- Take $a_t \sim \pi_b(s_t)$: Sample action from ϵ -greedy policy. Here π_b is ϵ -greedy w.r.t. Q .
- It builds results through stochastic approximation, decreasing step sizes to satisfy Bellman contraction property, and utilizing bounded rewards and value function.

Q-learning for finite-state and finite-action MDPs converges to optimal value $Q^*(s, a)$ given:

- Policy sequence $\pi_t(a|s)$ satisfies conditions of GLIE.

- $\sum_{t=1}^{\infty} \alpha_t = \infty$ and $\sum_{t=1}^{\infty} \alpha_t^2 < \infty$

3.1 Function Approximation

Function **approximation** carries the following two purposes:

1. Avoid [explicit storing/ learning](#) for every single state and action (model/ value/ policy)
2. [Compact representation](#), generalized across states and actions (i.e. less memory and computation needed)

There are **THREE** forms of approximation: [bootstrapping](#), [sampling](#), and [VFA](#).

[With a model](#), we carry out VFA by assuming we can "query (s, a) " to return $Q^{\pi}(s, a)$, i.e. supervised learning:

Model-based VFA: Process

- **Objective:** approximate representation of Q^π given parameterized form $\hat{Q}(s, a; \mathbf{w})$
- **Input:** \mathbf{w} is a parameter vector of weights.
- **Method:** Minimize loss function between $Q^\pi(s, a)$ and approximation $\hat{Q}(s, a; \mathbf{w})$:

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[\left(Q^\pi(s, a) - \hat{Q}(s, a; \mathbf{w}) \right)^2 \right]$$

- **Means:** Find local minimum with gradient descent

$$\delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$

- **Improvement:** Stochastic gradient descent (SGD) uses finite (often one) sample to approximate gradient

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = -2 \mathbb{E}_\pi [Q^\pi(s, a) - \hat{Q}(s, a; \mathbf{w})] \nabla_{\mathbf{w}} \hat{Q}$$

How does SGD differ from GD? What is "finite # samples"? [See answer from GPT]

3.1.1 VFA without Model

Without a model (oracle), follow policy π / prior data to estimate V^π and Q^π .

Model-free VFA: Process

- **Input:** Generated sample paths by following fixed policy π .
- **Method:** Through lookup table storing estimates of V^π or Q^π
→ Update estimates after each episode (MC)/ step (TD)
- Change the estimate update step to include fitting function approximator.
How? Return G_t is unbiased but noisy sample of true $Q^\pi(s_t, a_t)$
⇒ Do supervised learning on $\langle (s_i, a_i), G_i \rangle$.

3.2 VFA under TDL

Recall that TDL uses bootstrapping and sampling to approximate V^π :

$$V^\pi(s) = V^\pi(s) + \alpha (r + \gamma V^\pi(s') - V^\pi(s))$$

where the temporal target $r + \gamma V^\pi(s')$ is biased against $V^\pi(s)$.

- When with a look table, represent each state value with a separate table entry.
- in VFA, target is $r + \gamma V^\pi(s'; \mathbf{w})$, also biased against $V^\pi(s)$.
- Equivalent to supervised learning on $\langle (s_i, r_i), \gamma \hat{V}^\pi(s_{i+1}; \mathbf{w}) \rangle$ and finding weights \mathbf{w} to minimize MSE with SGD:

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[\left(r_j + \gamma V^\pi(s_{j+1}, \mathbf{w}) - \hat{V}(s_j; \mathbf{w}) \right)^2 \right]$$

Question: Why the function involve V here but not Q like in model-based VFA?

3.2.1 Control using VFA

Intuitively the few main points are:

- **Estimation:** $Q^\pi(s, a; \mathbf{w}) \approx Q^\pi$ by sampling gradients from SGD, for local minimum.

- **Interleaving:** Perform ϵ -greedy policy improvement between VFA approximations.

There are a few incremental approaches to determine the increment $\Delta \mathbf{w}$ for action-value function approximation:

1. **General form:** use **target value**

$$\Delta \mathbf{w} = \alpha(Q(s_t, a_t) - \hat{Q}(s_t, a_t; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_t, a_t; \mathbf{w})$$

2. **MC Simulation:** use **return** G_t as substitute target

$$\Delta \mathbf{w} = \alpha(G_t - \hat{Q}(s_t, a_t; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_t, a_t; \mathbf{w})$$

3. **SARSA:** use **TD target**, i.e. $r + \gamma \hat{Q}(s', a'; \mathbf{w})$

$$\Delta \mathbf{w} = \alpha(r + \gamma \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

4. **Q-Learning:** use **related TD target** to SARSA

$$\Delta \mathbf{w} = \alpha(r + \gamma \arg \max_{a'} Q(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

Q: What is the original purpose of interleaving? How does it cause problems? Q: Recall, how is the target value in general form VFA increment determined? Q: Recall why is there t in MC/ General form for s and a but NOT SARSA/ Q-learning? Q: Recall: What are the major differences between SARSA and Q-learning by tweaking the TD target?

3.3 Deadly Triad and Remedies

Applying bootstrapping, Function Approximation and Off-policy Learning leads to oscillations and non-convergence!

- Reason: Bellman operators on updates are contractions, but **VFA may be expansions**.
- Q-learning with VFA diverges with **correlation between samples** and **non-stationary targets**.

Solution presented includes **Deep Q-learning Network (DQN)**, featuring replay and fixed Q-targets.

Remedy 1: DQN with Replay

Extract **sample experience tuples** $(s_i, a, r, s_{i+1}) \sim D$ (a replay buffer from prior experience)

- Target value: $r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w})$; plug into SGD
- The replay buffer is usually of fixed size
- SGD to update network weights: $\Delta \mathbf{w} = \alpha \left(r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$

Remedy 2: Fix Target Weights

Fix the target weights used in the target calculation for multiple updates

- Different set of target network weights (\mathbf{w}^-) than the weights being updated (\mathbf{w})
- Fix \mathbf{w}^- for multiple updates before replacing it with new \mathbf{w} for next multiple steps
- Target value: $r + \gamma \max_{a'} \hat{Q}(s', a'; \mathbf{w}^-)$

3.4 DQN in Practice with Atari

Below discusses an example in practice:

Workflow of applying DQN in Atari

- Learn $Q(s, a)$ from pixel s ;
- Input state s = raw pixels from last 4 frames
- Output: $Q(s, a)$ for all (joystick, button)
- Reward: Δ Score
- Architecture: Deep NN with CNN

Outcome: DQN mostly outperforms best linear learners!

Caveats of Atari DQN

- Deep network *without* fixed Q in fact **hurts performance!** (worse than linear)
- Replays are hugely important!
- Immediate improvements: Double DQN (Van Hasselt, 2016), Prioritized Replay (Schaul et. al., 2016), Dueling DQN (Wang et. al., 2016)

Q: "Bellman operators are contractions, but value function approximation fitting can be an expansion". What does this mean? Q: How are replays hugely important beyond decorrelating between samples?